

```

1 import numpy as np
2 import keras
3 from keras import layers

```

▼ Prepare the data

```

1 # Model / data parameters
2 num_classes = 10
3 input_shape = (28, 28, 1)
4
5 # Load the data and split it between train and test sets
6 (x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
7
8

```

📄 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
 11490434/11490434 ————— 1s 0us/step

```

1 # Scale images to the [0, 1] range
2 x_train = x_train.astype("float32") / 255
3 x_test = x_test.astype("float32") / 255
4 # Make sure images have shape (28, 28, 1)
5 x_train = np.expand_dims(x_train, -1)
6 x_test = np.expand_dims(x_test, -1)
7 print("x_train shape:", x_train.shape)
8 print(x_train.shape[0], "train samples")
9 print(x_test.shape[0], "test samples")
10
11
12 # convert class vectors to binary class matrices
13 y_train = keras.utils.to_categorical(y_train, num_classes)
14 y_test = keras.utils.to_categorical(y_test, num_classes)

```

📄 x_train shape: (60000, 28, 28, 1)
 60000 train samples
 10000 test samples

▼ Build the model

```

1 model = keras.Sequential(
2     [
3         keras.Input(shape=input_shape),
4         layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
5         layers.MaxPooling2D(pool_size=(2, 2)),
6         layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
7         layers.MaxPooling2D(pool_size=(2, 2)),
8         layers.Flatten(),
9         layers.Dropout(0.5),
10        layers.Dense(num_classes, activation="softmax"),
11    ]
12 )
13
14 model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	
conv2d (Conv2D)	(None, 26, 26, 32)	
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	
conv2d_1 (Conv2D)	(None, 11, 11, 64)	
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	
flatten (Flatten)	(None, 1600)	
dropout (Dropout)	(None, 1600)	
dense (Dense)	(None, 10)	

Total params: 34,826 (136.04 KB)
 Trainable params: 34,826 (136.04 KB)
 Non-trainable params: 0 (0.00 KB)

Train the model

```
1 batch_size = 128
2 epochs = 15
3
4 model.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["accuracy"])
5
6 model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, validation_split=0.1)
7
```

Epoch 1/15
 422/422 ————— 57s 127ms/step - accuracy: 0.7690 - loss: 0.7500 - val_accuracy: 0.9755 - val_loss
 Epoch 2/15
 422/422 ————— 73s 107ms/step - accuracy: 0.9595 - loss: 0.1318 - val_accuracy: 0.9828 - val_loss
 Epoch 3/15
 422/422 ————— 43s 101ms/step - accuracy: 0.9711 - loss: 0.0928 - val_accuracy: 0.9860 - val_loss
 Epoch 4/15
 422/422 ————— 82s 101ms/step - accuracy: 0.9772 - loss: 0.0757 - val_accuracy: 0.9885 - val_loss
 Epoch 5/15
 422/422 ————— 42s 101ms/step - accuracy: 0.9804 - loss: 0.0641 - val_accuracy: 0.9895 - val_loss
 Epoch 6/15
 422/422 ————— 82s 101ms/step - accuracy: 0.9811 - loss: 0.0610 - val_accuracy: 0.9898 - val_loss
 Epoch 7/15
 422/422 ————— 44s 105ms/step - accuracy: 0.9824 - loss: 0.0536 - val_accuracy: 0.9893 - val_loss
 Epoch 8/15
 422/422 ————— 44s 103ms/step - accuracy: 0.9836 - loss: 0.0511 - val_accuracy: 0.9902 - val_loss
 Epoch 9/15
 422/422 ————— 82s 103ms/step - accuracy: 0.9851 - loss: 0.0489 - val_accuracy: 0.9917 - val_loss
 Epoch 10/15
 422/422 ————— 82s 102ms/step - accuracy: 0.9850 - loss: 0.0462 - val_accuracy: 0.9908 - val_loss
 Epoch 11/15
 422/422 ————— 85s 110ms/step - accuracy: 0.9861 - loss: 0.0440 - val_accuracy: 0.9913 - val_loss
 Epoch 12/15
 422/422 ————— 48s 113ms/step - accuracy: 0.9873 - loss: 0.0410 - val_accuracy: 0.9918 - val_loss
 Epoch 13/15
 422/422 ————— 77s 100ms/step - accuracy: 0.9879 - loss: 0.0362 - val_accuracy: 0.9923 - val_loss
 Epoch 14/15
 422/422 ————— 43s 101ms/step - accuracy: 0.9878 - loss: 0.0370 - val_accuracy: 0.9920 - val_loss
 Epoch 15/15
 422/422 ————— 82s 101ms/step - accuracy: 0.9888 - loss: 0.0341 - val_accuracy: 0.9915 - val_loss
 <keras.src.callbacks.history.History at 0x7f753a3936d0>

Evaluate the trained model

```
1 score = model.evaluate(x_test, y_test, verbose=0)
2 print("Test loss:", score[0])
3 print("Test accuracy:", score[1])
```

Test loss: 0.02719920687377453
 Test accuracy: 0.9908999800682068

1 Start coding or [generate](#) with AI.

