

✓ Copyright 2022 The TensorFlow Authors.

> Licensed under the Apache License, Version 2.0 (the "License");

[Show code](#)



[View on TensorFlow.org](#)



[Run in Google Colab](#)



[View source on GitHub](#)



[Download notebook](#)

✓ Video classification with a 3D convolutional neural network

This tutorial demonstrates training a 3D convolutional neural network (CNN) for video classification using the [UCF101](#) action recognition dataset. A 3D CNN uses a three-dimensional filter to perform convolutions. The kernel is able to slide in three directions, whereas in a 2D CNN it can slide in two dimensions. The model is based on the work published in [A Closer Look at Spatiotemporal Convolutions for Action Recognition](#) by D. Tran et al. (2017). In this tutorial, you will:

- Build an input pipeline
- Build a 3D convolutional neural network model with residual connections using Keras functional API
- Train the model
- Evaluate and test the model

This video classification tutorial is the second part in a series of TensorFlow video tutorials. Here are the other three tutorials:

- [Load video data](#): This tutorial explains much of the code used in this document.
- [MoViNet for streaming action recognition](#): Get familiar with the MoViNet models that are available on TF Hub.
- [Transfer learning for video classification with MoViNet](#): This tutorial explains how to use a pre-trained video classification model trained on a different dataset with the UCF-101 dataset.

✓ Setup

Begin by installing and importing some necessary libraries, including: [remotezip](#) to inspect the contents of a ZIP file, [tqdm](#) to use a progress bar, [OpenCV](#) to process video files, [einops](#) for

performing more complex tensor operations, and [tensorflow_docs](#) for embedding data in a Jupyter notebook.

```
1 !pip install remotezip tqdm opencv-python einops
2 !pip install -U tensorflow keras
```



Collecting remotezip

```
Downloading remotezip-0.12.3-py3-none-any.whl.metadata (7.2 kB)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: opencv-python in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: einops in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: numpy>=1.21.2 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages
Downloading remotezip-0.12.3-py3-none-any.whl (8.1 kB)
```

Installing collected packages: remotezip

Successfully installed remotezip-0.12.3

```
Requirement already satisfied: tensorflow in /usr/local/lib/python3.10/dist-packages
```

Collecting tensorflow

```
Downloading tensorflow-2.18.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (452.4 MB)
```

```
Requirement already satisfied: keras in /usr/local/lib/python3.10/dist-packages
```

Collecting keras

```
Downloading keras-3.7.0-py3-none-any.whl.metadata (5.8 kB)
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.10/dist-packages
Collecting tensorboard<2.19,>=2.18 (from tensorflow)
```

```
Downloading tensorboard-2.18.0-py3-none-any.whl.metadata (1.6 kB)
```

```
Requirement already satisfied: numpy<2.1.0,>=1.26.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: h5py>=3.11.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: ml-dtypes<0.5.0,>=0.4.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: rich in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: namex in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: optree in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.10/dist-packages
```

```
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.10/
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/pytho
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/pyt
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.1
```

```
1 import tqdm
2 import random
3 import pathlib
4 import itertools
5 import collections
6
7 import cv2
8 import einops
9 import numpy as np
10 import remotezip as rz
11 import seaborn as sns
12 import matplotlib.pyplot as plt
13
14 import tensorflow as tf
15 import keras
16 from keras import layers
```

✓ Load and preprocess video data

The hidden cell below defines helper functions to download a slice of data from the UCF-101 dataset, and load it into a `tf.data.Dataset`. You can learn more about the specific preprocessing steps in the [Loading video data tutorial](#), which walks you through this code in more detail.

The `FrameGenerator` class at the end of the hidden block is the most important utility here. It creates an iterable object that can feed data into the TensorFlow data pipeline. Specifically, this class contains a Python generator that loads the video frames along with its encoded label. The generator (`__call__`) function yields the frame array produced by `frames_from_video_file` and a one-hot encoded vector of the label associated with the set of frames.

```
1 #@title
2
3 def list_files_per_class(zip_url):
4     """
5     List the files in each class of the dataset given the zip URL.
6
7     Args:
8         zip_url: URL from which the files can be unzipped.
9
10    Return:
11        files: List of files in each of the classes.
12    """
13    files = []
14    with rz.RemoteZip(URL) as zip:
15        for zip_info in zip.infolist():
16            files.append(zip_info.filename)
17    return files
18
19 def get_class(fname):
20     """
21     Retrieve the name of the class given a filename.
22
23     Args:
24         fname: Name of the file in the UCF101 dataset.
25
26     Return:
27         Class that the file belongs to.
28     """
29    return fname.split('_')[-3]
30
31 def get_files_per_class(files):
32     """
33     Retrieve the files that belong to each class.
34
35     Args:
36         files: List of files in the dataset.
37
38     Return:
39         Dictionary of class names (key) and files (values).
40     """
41    files_for_class = collections.defaultdict(list)
42    for fname in files:
43        class_name = get_class(fname)
44        files_for_class[class_name].append(fname)
45    return files_for_class
46
47 def download_from_zip(zip_url, to_dir, file_names):
48     """
49     Download the contents of the zip file from the zip URL.
50
51     Args:
52         zip_url: Zip URL containing data.
53         to_dir: Directory to download data to.
54         file_names: Names of files to download.
55     """
```



```

56 with rz.RemoteZip(zip_url) as zip:
57     for fn in tqdm.tqdm(file_names):
58         class_name = get_class(fn)
59         zip.extract(fn, str(to_dir / class_name))
60         unzipped_file = to_dir / class_name / fn
61
62         fn = pathlib.Path(fn).parts[-1]
63         output_file = to_dir / class_name / fn
64         unzipped_file.rename(output_file,)
65
66 def split_class_lists(files_for_class, count):
67     """
68     Returns the list of files belonging to a subset of data as well as the rem
69     files that need to be downloaded.
70
71     Args:
72         files_for_class: Files belonging to a particular class of data.
73         count: Number of files to download.
74
75     Return:
76         split_files: Files belonging to the subset of data.
77         remainder: Dictionary of the remainder of files that need to be download
78     """
79     split_files = []
80     remainder = {}
81     for cls in files_for_class:
82         split_files.extend(files_for_class[cls][:count])
83         remainder[cls] = files_for_class[cls][count:]
84     return split_files, remainder
85
86 def download UFC101_subset(zip_url, num_classes, splits, download_dir):
87     """
88     Download a subset of the UFC101 dataset and split them into various parts,
89     training, validation, and test.
90
91     Args:
92         zip_url: Zip URL containing data.
93         num_classes: Number of labels.
94         splits: Dictionary specifying the training, validation, test, etc. (key)
95                 (value is number of files per split).
96         download_dir: Directory to download data to.
97
98     Return:
99         dir: Posix path of the resulting directories containing the splits of da
100     """
101     files = list_files_per_class(zip_url)
102     for f in files:
103         tokens = f.split('/')
104         if len(tokens) <= 2:
105             files.remove(f) # Remove that item from the list if it does not have a f
106
107     files_for_class = get_files_per_class(files)
108
109     classes = list(files_for_class.keys())[:num_classes]
110

```



```
111 for cls in classes:
112     new_files_for_class = files_for_class[cls]
113     random.shuffle(new_files_for_class)
114     files_for_class[cls] = new_files_for_class
115
116 # Only use the number of classes you want in the dictionary
117 files_for_class = {x: files_for_class[x] for x in list(files_for_class)[:num]}
118
119 dirs = {}
120 for split_name, split_count in splits.items():
121     print(split_name, ":")
122     split_dir = download_dir / split_name
123     split_files, files_for_class = split_class_lists(files_for_class, split_count)
124     download_from_zip(zip_url, split_dir, split_files)
125     dirs[split_name] = split_dir
126
127 return dirs
128
129 def format_frames(frame, output_size):
130     """
131     Pad and resize an image from a video.
132
133     Args:
134         frame: Image that needs to be resized and padded.
135         output_size: Pixel size of the output frame image.
136
137     Return:
138         Formatted frame with padding of specified output size.
139     """
140     frame = tf.image.convert_image_dtype(frame, tf.float32)
141     frame = tf.image.resize_with_pad(frame, *output_size)
142     return frame
143
144 def frames_from_video_file(video_path, n_frames, output_size = (224,224), frame_step = 1):
145     """
146     Creates frames from each video file present for each category.
147
148     Args:
149         video_path: File path to the video.
150         n_frames: Number of frames to be created per video file.
151         output_size: Pixel size of the output frame image.
152
153     Return:
154         An NumPy array of frames in the shape of (n_frames, height, width, channels)
155     """
156     # Read each video frame by frame
157     result = []
158     src = cv2.VideoCapture(str(video_path))
159
160     video_length = src.get(cv2.CAP_PROP_FRAME_COUNT)
161
162     need_length = 1 + (n_frames - 1) * frame_step
163
164     if need_length > video_length:
165         start = 0
```



```

166 else:
167     max_start = video_length - need_length
168     start = random.randint(0, max_start + 1)
169
170 src.set(cv2.CAP_PROP_POS_FRAMES, start)
171 # ret is a boolean indicating whether read was successful, frame is the image
172 ret, frame = src.read()
173 result.append(format_frames(frame, output_size))
174
175 for _ in range(n_frames - 1):
176     for _ in range(frame_step):
177         ret, frame = src.read()
178         if ret:
179             frame = format_frames(frame, output_size)
180             result.append(frame)
181         else:
182             result.append(np.zeros_like(result[0]))
183 src.release()
184 result = np.array(result)[..., [2, 1, 0]]
185
186 return result
187
188 class FrameGenerator:
189     def __init__(self, path, n_frames, training = False):
190         """ Returns a set of frames with their associated label.
191
192         Args:
193             path: Video file paths.
194             n_frames: Number of frames.
195             training: Boolean to determine if training dataset is being created.
196         """
197         self.path = path
198         self.n_frames = n_frames
199         self.training = training
200         self.class_names = sorted(set(p.name for p in self.path.iterdir() if p.is_file()))
201         self.class_ids_for_name = dict((name, idx) for idx, name in enumerate(self.class_names))
202
203     def get_files_and_class_names(self):
204         video_paths = list(self.path.glob('*/*.avi'))
205         classes = [p.parent.name for p in video_paths]
206         return video_paths, classes
207
208     def __call__(self):
209         video_paths, classes = self.get_files_and_class_names()
210
211         pairs = list(zip(video_paths, classes))
212
213         if self.training:
214             random.shuffle(pairs)
215
216         for path, name in pairs:
217             video_frames = frames_from_video_file(path, self.n_frames)
218             label = self.class_ids_for_name[name] # Encode labels
219             yield video_frames, label

```

```

1 URL = 'https://storage.googleapis.com/thumos14_files/UCF101_videos.zip'
2 download_dir = pathlib.Path('./UCF101_subset/')
3 subset_paths = download_ufc_101_subset(URL,
4                                     num_classes = 10,
5                                     splits = {"train": 30, "val": 10, "test": 10},
6                                     download_dir = download_dir)

```

```

⇒ train :
100%|██████████| 300/300 [02:38<00:00, 1.89it/s]
val :
100%|██████████| 100/100 [00:52<00:00, 1.91it/s]
test :
100%|██████████| 100/100 [00:52<00:00, 1.90it/s]

```

Create the training, validation, and test sets (train_ds, val_ds, and test_ds).

```

1 n_frames = 10
2 batch_size = 8
3
4 output_signature = (tf.TensorSpec(shape = (None, None, None, 3), dtype = tf.float32),
5                     tf.TensorSpec(shape = (), dtype = tf.int16))
6
7 train_ds = tf.data.Dataset.from_generator(FrameGenerator(subset_paths['train'],
8                                                         output_signature = output_signature))
9
10
11 # Batch the data
12 train_ds = train_ds.batch(batch_size)
13
14 val_ds = tf.data.Dataset.from_generator(FrameGenerator(subset_paths['val'], n_frames,
15                                                         output_signature = output_signature))
16 val_ds = val_ds.batch(batch_size)
17
18 test_ds = tf.data.Dataset.from_generator(FrameGenerator(subset_paths['test'], n_frames,
19                                                         output_signature = output_signature))
20
21 test_ds = test_ds.batch(batch_size)

```


✓ Create the model

The following 3D convolutional neural network model is based off the paper [A Closer Look at Spatiotemporal Convolutions for Action Recognition](#) by D. Tran et al. (2017). The paper compares several versions of 3D ResNets. Instead of operating on a single image with dimensions (height, width), like standard ResNets, these operate on video volume (time, height, width). The most obvious approach to this problem would be replace each 2D convolution (layers.Conv2D) with a 3D convolution (layers.Conv3D).

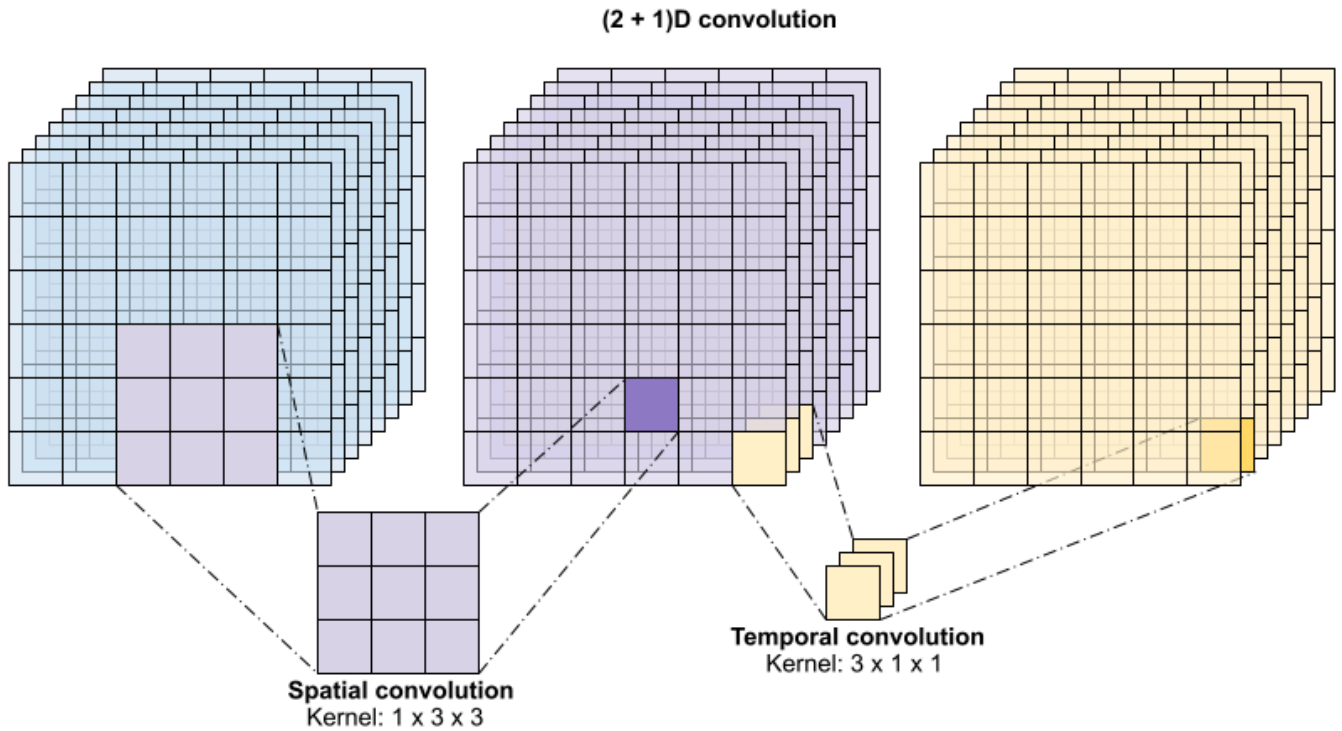
This tutorial uses a (2 + 1)D convolution with [residual connections](#). The (2 + 1)D convolution allows for the decomposition of the spatial and temporal dimensions, therefore creating two

separate steps. An advantage of this approach is that factorizing the convolutions into spatial and temporal dimensions saves parameters.

For each output location a 3D convolution combines all the vectors from a 3D patch of the volume to create one vector in the output volume.



This operation takes $\text{time} \times \text{height} \times \text{width} \times \text{channels}$ inputs and produces channel outputs (assuming the number of input and output channels are the same). So a 3D convolution layer with a kernel size of $(3 \times 3 \times 3)$ would need a weight-matrix with $27 \times \text{channels} \times 2$ entries. The reference paper found that a more effective & efficient approach was to factorize the convolution. Instead of a single 3D convolution to process the time and space dimensions, they proposed a "(2+1)D" convolution which processes the space and time dimensions separately. The figure below shows the factored spatial and temporal convolutions of a $(2 + 1)$ D convolution.



The main advantage of this approach is that it reduces the number of parameters. In the (2 + 1)D convolution the spatial convolution takes in data of the shape (1, width, height), while the temporal convolution takes in data of the shape (time, 1, 1). For example, a (2 + 1)D convolution with kernel size (3 x 3 x 3) would need weight matrices of size (9 * channels**2) + (3 * channels**2), less than half as many as the full 3D convolution. The tutorial implements (2 + 1)D ResNet18, where each convolution in the resnet is replaced by a (2+1)D convolution.

```
1 # Define the dimensions of one frame in the set of frames created
2 HEIGHT = 224
3 WIDTH = 224
```

```

1 class Conv2Plus1D(keras.layers.Layer):
2     def __init__(self, filters, kernel_size, padding):
3         """
4         A sequence of convolutional layers that first apply the convolution oper
5         spatial dimensions, and then the temporal dimension.
6         """
7         super().__init__()
8         self.seq = keras.Sequential([
9             # Spatial decomposition
10            layers.Conv3D(filters=filters,
11                           kernel_size=(1, kernel_size[1], kernel_size[2]),
12                           padding=padding),
13            # Temporal decomposition
14            layers.Conv3D(filters=filters,
15                           kernel_size=(kernel_size[0], 1, 1),
16                           padding=padding)
17        ])
18
19     def call(self, x):
20         return self.seq(x)

```

A ResNet model is made from a sequence of residual blocks. A residual block has two branches. The main branch performs the calculation, but is difficult for gradients to flow through. The residual branch bypasses the main calculation and mostly just adds the input to the output of the main branch. Gradients flow easily through this branch. Therefore, an easy path from the loss function to any of the residual block's main branch will be present. This avoids the vanishing gradient problem.

Create the main branch of the residual block with the following class. In contrast to the standard ResNet structure this uses the custom `Conv2Plus1D` layer instead of `layers.Conv2D`.

```

1 class ResidualMain(keras.layers.Layer):
2     """
3     Residual block of the model with convolution, layer normalization, and the
4     activation function, ReLU.
5     """
6     def __init__(self, filters, kernel_size):
7         super().__init__()
8         self.seq = keras.Sequential([
9             Conv2Plus1D(filters=filters,
10                        kernel_size=kernel_size,
11                        padding='same'),
12             layers.LayerNormalization(),
13             layers.ReLU(),
14             Conv2Plus1D(filters=filters,
15                        kernel_size=kernel_size,
16                        padding='same'),
17             layers.LayerNormalization()
18         ])
19
20     def call(self, x):
21         return self.seq(x)

```

To add the residual branch to the main branch it needs to have the same size. The Project layer below deals with cases where the number of channels is changed on the branch. In particular, a sequence of densely-connected layer followed by normalization is added.

```

1 class Project(keras.layers.Layer):
2     """
3     Project certain dimensions of the tensor as the data is passed through dif
4     sized filters and downsampled.
5     """
6     def __init__(self, units):
7         super().__init__()
8         self.seq = keras.Sequential([
9             layers.Dense(units),
10            layers.LayerNormalization()
11        ])
12
13     def call(self, x):
14         return self.seq(x)

```

Use `add_residual_block` to introduce a skip connection between the layers of the model.

```

1 def add_residual_block(input, filters, kernel_size):
2     """
3     Add residual blocks to the model. If the last dimensions of the input data
4     and filter size does not match, project it such that last dimension matche
5     """
6     out = ResidualMain(filters,
7                         kernel_size)(input)
8
9     res = input
10    # Using the Keras functional APIs, project the last dimension of the tensor
11    # match the new filter size
12    if out.shape[-1] != input.shape[-1]:
13        res = Project(out.shape[-1])(res)
14
15    return layers.add([res, out])

```

Resizing the video is necessary to perform downsampling of the data. In particular, downsampling the video frames allow for the model to examine specific parts of frames to detect patterns that may be specific to a certain action. Through downsampling, non-essential information can be discarded. Moreover, resizing the video will allow for dimensionality reduction and therefore faster processing through the model.

```

1 class ResizeVideo(keras.layers.Layer):
2     def __init__(self, height, width):
3         super().__init__()
4         self.height = height
5         self.width = width
6         self.resizing_layer = layers.Resizing(self.height, self.width)
7
8     def call(self, video):
9         """
10        Use the einops library to resize the tensor.
11
12        Args:
13            video: Tensor representation of the video, in the form of a set of fra
14
15        Return:
16            A downsampled size of the video according to the new height and width
17        """
18        # b stands for batch size, t stands for time, h stands for height,
19        # w stands for width, and c stands for the number of channels.
20        old_shape = einops.parse_shape(video, 'b t h w c')
21        images = einops.rearrange(video, 'b t h w c -> (b t) h w c')
22        images = self.resizing_layer(images)
23        videos = einops.rearrange(
24            images, '(b t) h w c -> b t h w c',
25            t = old_shape['t'])
26        return videos

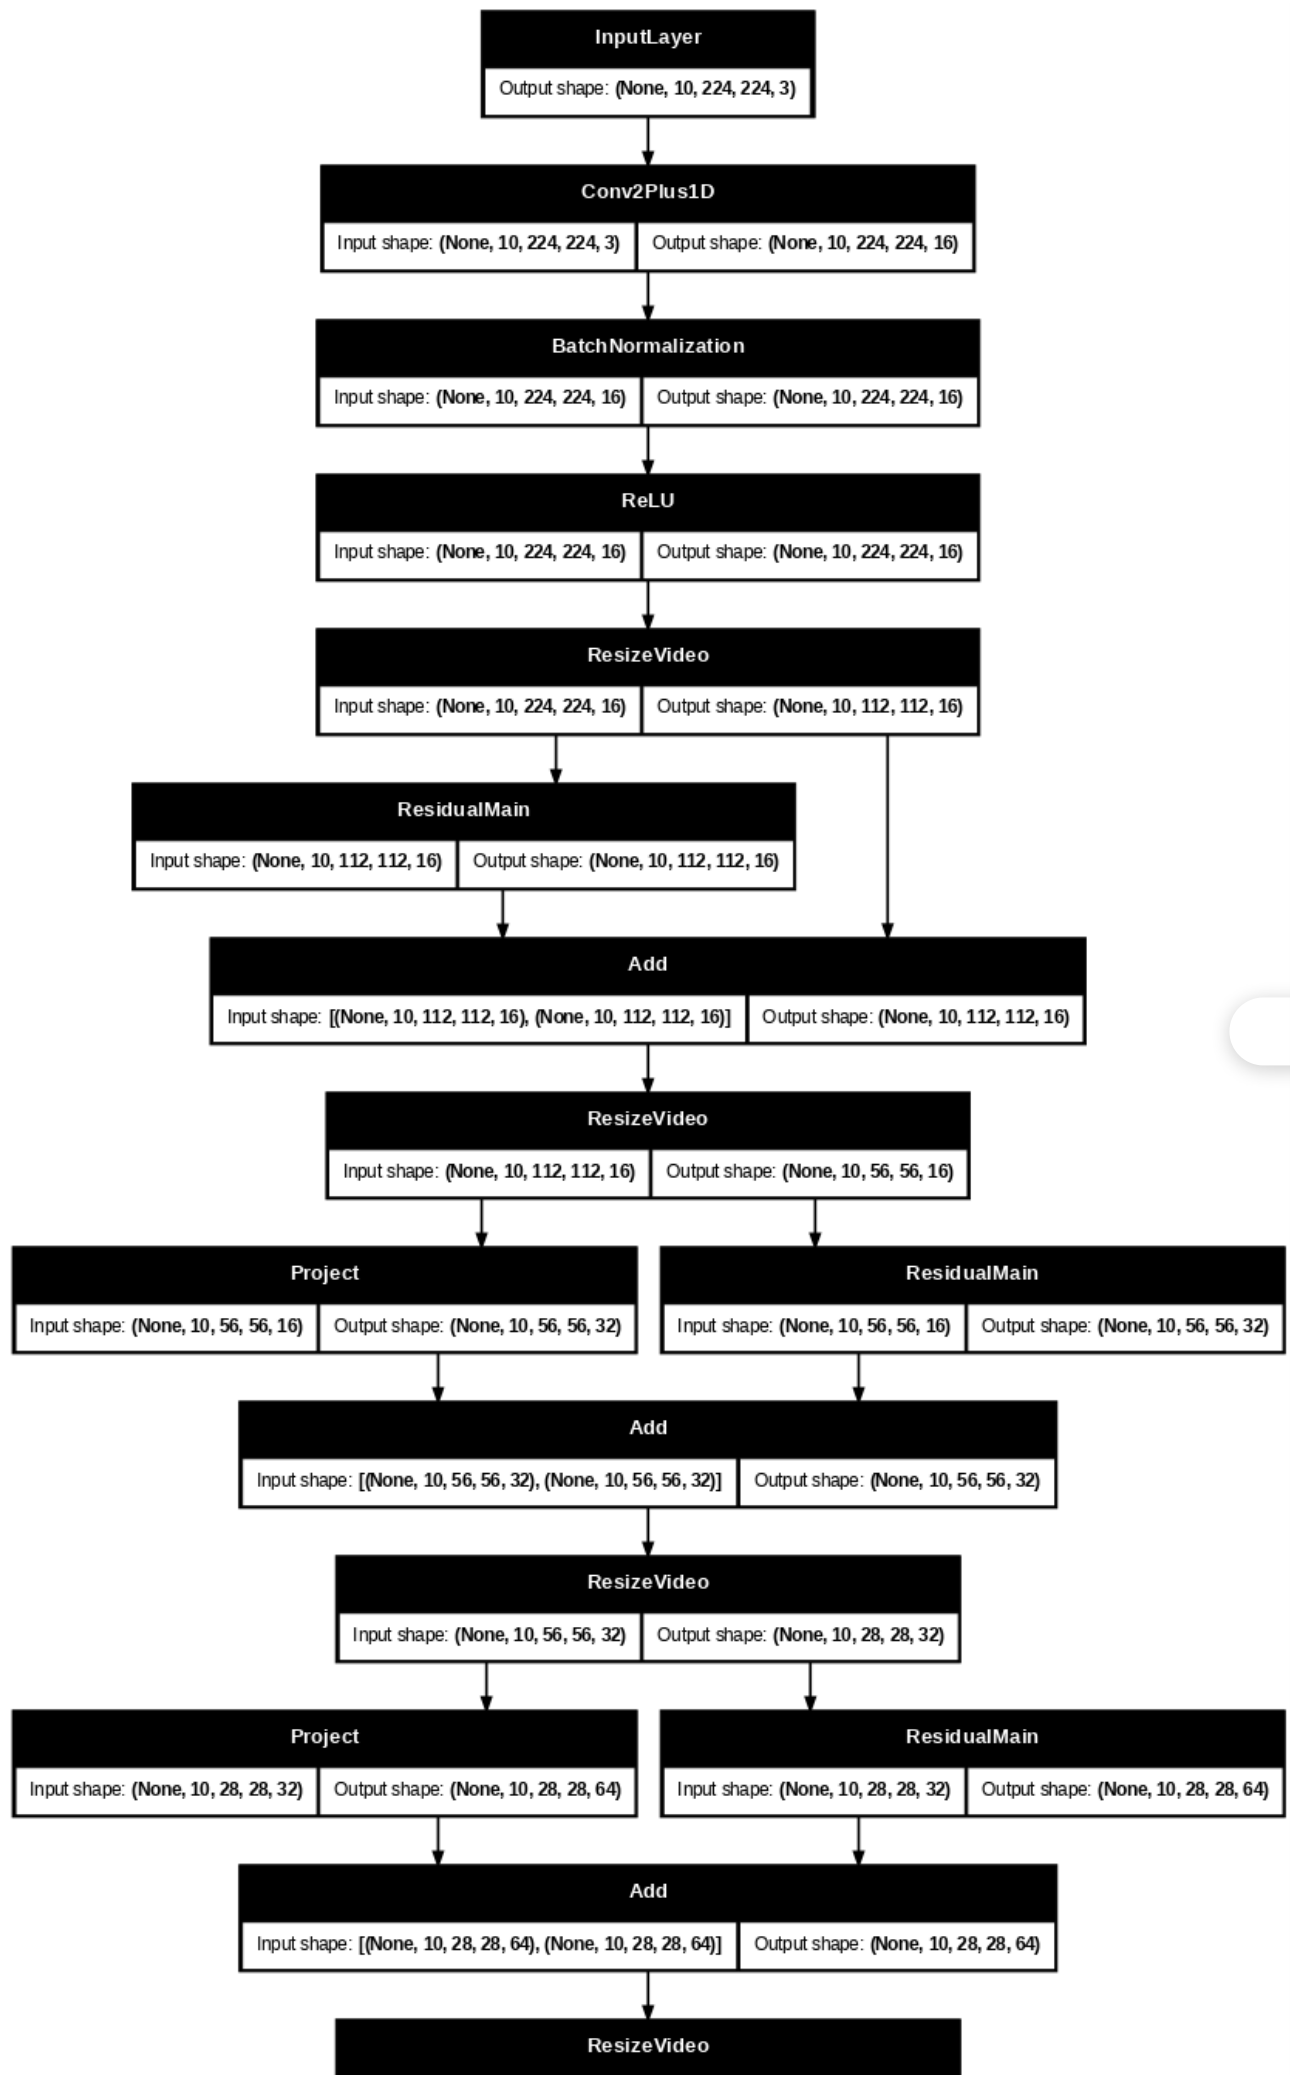
```

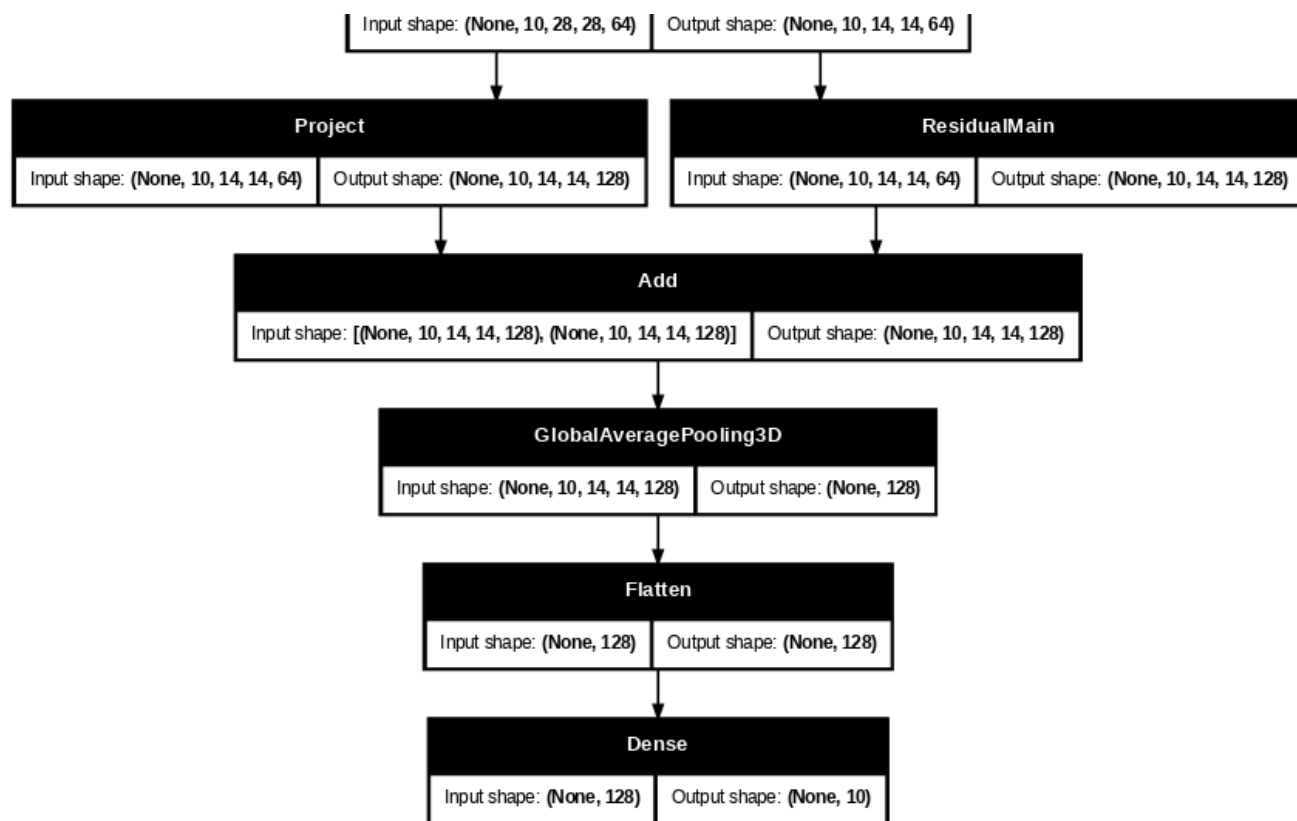
Use the [Keras functional API](#) to build the residual network.

```
1 input_shape = (None, 10, HEIGHT, WIDTH, 3)
2 input = layers.Input(shape=(input_shape[1:]))
3 x = input
4
5 x = Conv2Plus1D(filters=16, kernel_size=(3, 7, 7), padding='same')(x)
6 x = layers.BatchNormalization()(x)
7 x = layers.ReLU()(x)
8 x = ResizeVideo(HEIGHT // 2, WIDTH // 2)(x)
9
10 # Block 1
11 x = add_residual_block(x, 16, (3, 3, 3))
12 x = ResizeVideo(HEIGHT // 4, WIDTH // 4)(x)
13
14 # Block 2
15 x = add_residual_block(x, 32, (3, 3, 3))
16 x = ResizeVideo(HEIGHT // 8, WIDTH // 8)(x)
17
18 # Block 3
19 x = add_residual_block(x, 64, (3, 3, 3))
20 x = ResizeVideo(HEIGHT // 16, WIDTH // 16)(x)
21
22 # Block 4
23 x = add_residual_block(x, 128, (3, 3, 3))
24
25 x = layers.GlobalAveragePooling3D()(x)
26 x = layers.Flatten()(x)
27 x = layers.Dense(10)(x)
28
29 model = keras.Model(input, x)

1 frames, label = next(iter(train_ds))
2 model.build(frames)

1 # Visualize the model
2 keras.utils.plot_model(model, expand_nested=True, dpi=60, show_shapes=True)
```



✓ Train the model

For this tutorial, choose the `tf.keras.optimizers.Adam` optimizer and the `tf.keras.losses.SparseCategoricalCrossentropy` loss function. Use the `metrics` argument to view the accuracy of the model performance at every step.

```
1 model.compile(loss = keras.losses.SparseCategoricalCrossentropy(from_logits=True),
2               optimizer = keras.optimizers.Adam(learning_rate = 0.0001),
3               metrics = ['accuracy'])
```

Train the model for 50 epochs with the Keras `Model.fit` method.

Note: This example model is trained on fewer data points (300 training and 100 validation examples) to keep training time reasonable for this tutorial. Moreover, this example model may take over one hour to train.

```
1 history = model.fit(x = train_ds,
2                    epochs = 50,
3                    validation_data = val_ds)
```



```
38/38 ————— 45/s 12s/step - accuracy: 0.7793 - loss: 0.6159 ·  
Epoch 44/50  
38/38 ————— 443s 12s/step - accuracy: 0.7771 - loss: 0.6183 ·  
Epoch 45/50  
38/38 ————— 464s 12s/step - accuracy: 0.8065 - loss: 0.5311 ·  
Epoch 46/50  
38/38 ————— 460s 12s/step - accuracy: 0.8445 - loss: 0.5099 ·  
Epoch 47/50  
38/38 ————— 442s 12s/step - accuracy: 0.8427 - loss: 0.4836 ·  
Epoch 48/50  
38/38 ————— 460s 12s/step - accuracy: 0.8750 - loss: 0.4656 ·  
Epoch 49/50  
38/38 ————— 505s 12s/step - accuracy: 0.7711 - loss: 0.6210 ·  
Epoch 50/50  
38/38 ————— 502s 12s/step - accuracy: 0.8838 - loss: 0.4403 ·
```

✓ Visualize the results

Create plots of the loss and accuracy on the training and validation sets:

```
1 def plot_history(history):
2     """
3     Plotting training and validation learning curves.
4
5     Args:
6         history: model history with all the metric measures
7     """
8     fig, (ax1, ax2) = plt.subplots(2)
9
10    fig.set_size_inches(18.5, 10.5)
11
12    # Plot loss
13    ax1.set_title('Loss')
14    ax1.plot(history.history['loss'], label = 'train')
15    ax1.plot(history.history['val_loss'], label = 'test')
16    ax1.set_ylabel('Loss')
17
18    # Determine upper bound of y-axis
19    max_loss = max(history.history['loss'] + history.history['val_loss'])
20
21    ax1.set_ylim([0, np.ceil(max_loss)])
22    ax1.set_xlabel('Epoch')
23    ax1.legend(['Train', 'Validation'])
24
25    # Plot accuracy
26    ax2.set_title('Accuracy')
27    ax2.plot(history.history['accuracy'], label = 'train')
28    ax2.plot(history.history['val_accuracy'], label = 'test')
29    ax2.set_ylabel('Accuracy')
30    ax2.set_ylim([0, 1])
31    ax2.set_xlabel('Epoch')
32    ax2.legend(['Train', 'Validation'])
33
34    plt.show()
35
36 plot_history(history)
```

✓ Evaluate the model

Use `Keras Model.evaluate` to get the loss and accuracy on the test dataset.

Note: The example model in this tutorial uses a subset of the UCF101 dataset to keep training time reasonable. The accuracy and loss can be improved with further hyperparameter tuning or more training data.

Accuracy

```
1 model.evaluate(test_ds, return_dict=True)
```

13/13 ————— 42s 3s/step - accuracy: 0.7599 - loss: 0.7216
 {'accuracy': 0.7300000190734863, 'loss': 0.8357596397399902}

0.4 |

To visualize model performance further, use a [confusion matrix](#). The confusion matrix allows you to assess the performance of the classification model beyond accuracy. In order to build the confusion matrix for this multi-class classification problem, get the actual values in the test set and the predicted values.

```
1 def get_actual_predicted_labels(dataset):
2     """
3     Create a list of actual ground truth values and the predictions from the
4
5     Args:
6         dataset: An iterable data structure, such as a TensorFlow Dataset, with
7
8     Return:
9         Ground truth and predicted values for a particular dataset.
10    """
11    actual = [labels for _, labels in dataset.unbatch()]
12    predicted = model.predict(dataset)
13
14    actual = tf.stack(actual, axis=0)
15    predicted = tf.concat(predicted, axis=0)
16    predicted = tf.argmax(predicted, axis=1)
17
18    return actual, predicted
```